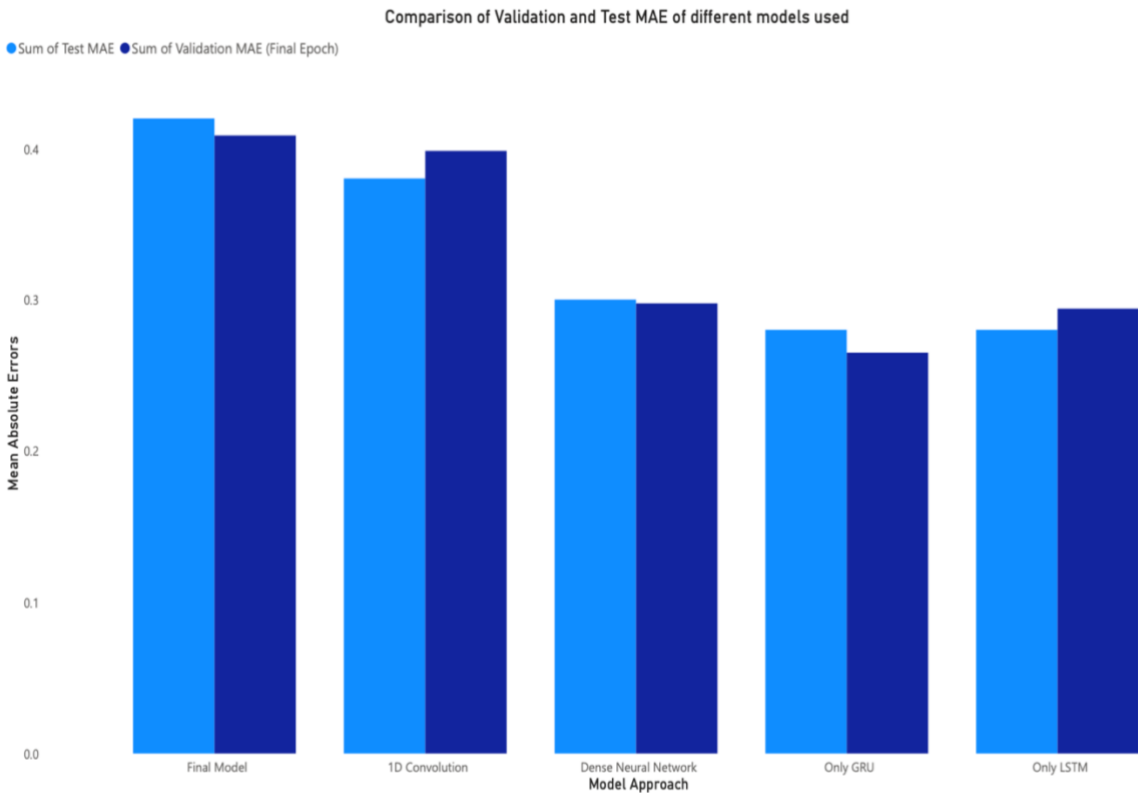


The project utilizes machine learning models to predict temperature based on historical data from the Jena Climate dataset spanning from 2009 to 2016. The project explored various modeling techniques, including dense neural networks, 1D convolutions, LSTM (Long Short-Term Memory), and GRU (Gated Recurrent Unit) layers, with the aim to minimize the mean absolute error (MAE) in temperature predictions.

Here's a table summarizing the key results from the different modeling approaches used in the weather forecasting study:

Model Approach	Training MAE (Final Epoch)	Validation MAE (Final Epoch)	Test MAE
Dense Neural Network	0.2314	0.2975	0.30
1D Convolution	0.2769	0.3986	0.38
Only LSTM	0.2262	0.2940	0.28
Only GRU	0.2883	0.2649	0.28
Final Model	0.3807	0.4088	0.42



Key Findings:

Dense Neural Network: Initial attempts with a simple dense network yielded a test MAE of 0.30.

1D Convolution: Employing a 1D convolution method resulted in a slightly higher test MAE of 0.38, suggesting that this approach may not be as effective for the given time-series data.

Only LSTM: The application of LSTM models showed promising results, reducing the test MAE to 0.28, indicating a better capability to capture temporal dependencies.

Only GRU: Introduction of GRU layers and dropout to combat overfitting achieved a test MAE of 0.28, consistent with LSTM performance but with a simpler architecture.

Final Model Construction:

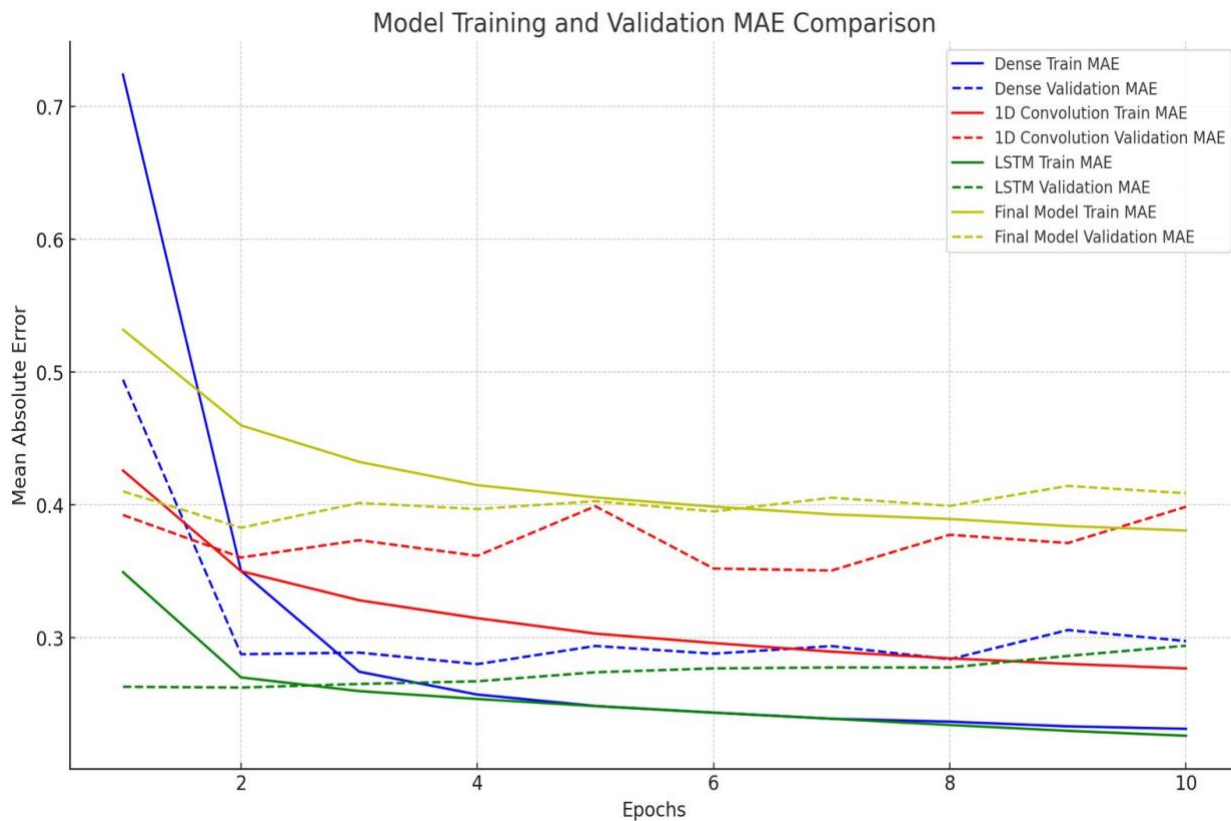
The ultimate approach combined LSTM layers with increased units, recurrent dropout, 1D convolutions, and dropout layers in a stacked model setup. Despite the complexity, this model achieved a slightly higher test MAE of 0.42, possibly due to overfitting or the added model complexity not translating to better generalization.

Conclusion:

The exploration of different neural network architectures to forecast weather demonstrated that LSTM and GRU models, with their ability to capture sequential data dependencies, were more effective than traditional dense networks and 1D convolutions. However, the final complex model combining LSTM and 1D convolutions did not significantly outperform simpler LSTM or GRU models, highlighting the importance of model simplicity and the diminishing returns of increased complexity.

Visualization

We visualized the performance of each model approach by plotting their mean absolute error (MAE) over the epochs for both training and validation datasets, to better illustrate the learning process and model performance over time.



The graph above illustrates the mean absolute error (MAE) for both training and validation sets across different modeling approaches used in the study: Dense Neural Network, 1D Convolution, LSTM, and the final complex model.

Dense Neural Network and 1D Convolution approaches show a noticeable gap between training and validation MAE, suggesting a degree of overfitting, especially in the dense network model.

The 1D Convolution's validation MAE tends to be higher and less stable, indicating that it might not be the best approach for this time-series data.

LSTM models demonstrate a closer alignment between training and validation MAE, showcasing better generalization capabilities. The LSTM approach also achieves lower MAE values compared to the dense and 1D convolution models, highlighting its suitability for time-series forecasting.

The Final Model, while more complex, does not significantly outperform the simpler LSTM model in terms of MAE. The training and validation MAE are higher, which might be attributed to the complexity of the model potentially leading to overfitting or not capturing the temporal dependencies as effectively as expected.

In conclusion, the LSTM model stands out as the most effective approach for this specific weather forecasting problem, achieving the best balance between accuracy and generalization. The exploration of model complexity with the final model illustrates that more layers and parameters do not necessarily equate to better performance on unseen data, underscoring the importance of model selection and tuning in machine learning projects.