

# Weather Forecasting

April 8, 2024

## 1 Assignment : 2

### 1.1 Author : Dev , Sakshi

Importing all the modules

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn import preprocessing
```

Reading the csv file

```
[2]: df = pd.read_csv("/Users/devmarwah/Documents/MSBA assignments/Advanced Machine_
Learning/jena_climate_2009_2016.csv")
```

### Data Exploration

Displaying the head of our data to have a look at it :

```
[3]: df.head()
```

```
[3]:
```

	Date Time	p (mbar)	T (degC)	Tpot (K)	Tdew (degC)	rh (%)	\
0	01.01.2009 00:10:00	996.52	-8.02	265.40	-8.90	93.3	
1	01.01.2009 00:20:00	996.57	-8.41	265.01	-9.28	93.4	
2	01.01.2009 00:30:00	996.53	-8.51	264.91	-9.31	93.9	
3	01.01.2009 00:40:00	996.51	-8.31	265.12	-9.07	94.2	
4	01.01.2009 00:50:00	996.51	-8.27	265.15	-9.04	94.1	

	VPmax (mbar)	VPact (mbar)	VPdef (mbar)	sh (g/kg)	H2OC (mmol/mol)	\
0	3.33	3.11	0.22	1.94	3.12	
1	3.23	3.02	0.21	1.89	3.03	
2	3.21	3.01	0.20	1.88	3.02	
3	3.26	3.07	0.19	1.92	3.08	
4	3.27	3.08	0.19	1.92	3.09	

	rho (g/m**3)	wv (m/s)	max. wv (m/s)	wd (deg)
0	1307.75	1.03	1.75	152.3
1	1309.80	0.72	1.50	136.1
2	1310.24	0.19	0.63	171.6

3	1309.19	0.34	0.50	198.0
4	1309.00	0.32	0.63	214.3

Checking shape of our data


```
[4]: df.shape
```

```
[4]: (420451, 15)
```

Plotting temperature

```
[5]: plt.plot(range(420451),df.iloc[:,2])
```


```
[5]: [<matplotlib.lines.Line2D at 0x144aba190>]
```



Plotting the temperature data for only first 10 days :

```
[6]: plt.plot(range(1440),df.iloc[0:1440,2])
```

```
[6]: [<matplotlib.lines.Line2D at 0x15042ce50>]
```



Printing the size of training, validation and test samples that we will be using :

```
[7]: n_train = int(0.5*len(df))
      n_val = int(0.25*len(df))
      n_test = int(n_train-n_val)
      print("Train samples : ",n_train)
      print("Validation samples : ",n_val)
      print("Test samples : ",n_test)
```

```
Train samples : 210225
Validation samples : 105112
Test samples : 105113
```

### Data Preparation

```
[8]: # Storing standard deviation and mean for further use and normalizing data:
      dfs = df.drop('Date Time',axis=1).to_numpy()
      mean = dfs[:n_train].mean(axis=0)
      dfs -=mean
      std = dfs[:n_train].std(axis=0)
      dfs /= std
      print(std[1])
      print(mean[1])
```

```
8.770983608349352
8.825903294089667
```

```
[9]: temperature = dfs[:,1]b
```

## Model Construction

Diving data into training, validation and test dataset

```
[10]: sampling_rate = 6
sequence_length = 120
delay = sampling_rate * (sequence_length + 24 - 1)
batch_size = 256
```

Converting data frame to array, discarding date-time and converting values to float

```
[16]: dfs=dfs.astype('float32')
temperature = temperature.astype('float32')
```

```
[23]: from tensorflow import keras
Train = keras.utils.timeseries_dataset_from_array(
    dfs[:-delay],
    targets = temperature[delay:],
    sampling_rate=sampling_rate,
    sequence_length = sequence_length,
    batch_size=batch_size,
    start_index=0,
    shuffle=True,
    end_index=n_train
)
Validation = keras.utils.timeseries_dataset_from_array(
    dfs[:-delay],
    targets = temperature[delay:],
    sampling_rate=sampling_rate,
    sequence_length = sequence_length,
    batch_size=batch_size,
    start_index=n_train,
    shuffle=True,
    end_index=n_train+n_val
)
Test = keras.utils.timeseries_dataset_from_array(
    dfs[:-delay],
    targets = temperature[delay:],
    sampling_rate=sampling_rate,
    sequence_length = sequence_length,
    batch_size=batch_size,
    start_index=n_train+n_val,
    shuffle=True
)
```

Inspecting the output of our Train dataset :

```
[24]: for samples,targets in Train :  
      print("Sample shape : ",samples.shape)  
      print("Target shape :",targets.shape)  
      break
```

Sample shape : (256, 120, 14)

Target shape : (256,)

Making a simple dense network model to check performance :

```
[47]: from tensorflow import keras  
      from tensorflow.keras import layers  
      inputs = keras.Input(shape=(sequence_length, dfs.shape[-1]))  
      x = layers.Reshape((sequence_length * dfs.shape[-1],))(inputs)  
      x = layers.Flatten()(x)  
      x = layers.Dense(16, activation="relu")(x)  
      outputs = layers.Dense(1)(x)  
      model = keras.Model(inputs, outputs)  
  
      callbacks = [  
          keras.callbacks.ModelCheckpoint("/Users/devmarwah/Documents/MSBA_↵  
      ↵assignments/Advanced Machine Learning/Weather Forecasting/jena_dense.keras",  
          save_best_only=True)  
      ]  
  
      model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])  
      history = model.fit(Train,  
                          epochs=10,  
                          validation_data=Validation,  
                          callbacks=callbacks)  
      model = keras.models.load_model("/Users/devmarwah/Documents/MSBA assignments/  
      ↵Advanced Machine Learning/Weather Forecasting/jena_dense.keras")  
      print(f"Test MAE: {model.evaluate(Test)[1]:.2f}")
```

Epoch 1/10

819/819 8s 9ms/step -

loss: 1.4598 - mae: 0.7240 - val\_loss: 0.3780 - val\_mae: 0.4943

Epoch 2/10

819/819 8s 10ms/step -

loss: 0.2071 - mae: 0.3503 - val\_loss: 0.1352 - val\_mae: 0.2876

Epoch 3/10

819/819 8s 9ms/step -

loss: 0.1219 - mae: 0.2744 - val\_loss: 0.1343 - val\_mae: 0.2888

Epoch 4/10

819/819 8s 9ms/step -

loss: 0.1071 - mae: 0.2572 - val\_loss: 0.1276 - val\_mae: 0.2801


Epoch 5/10

```
819/819          8s 9ms/step -
loss: 0.1000 - mae: 0.2486 - val_loss: 0.1386 - val_mae: 0.2938
Epoch 6/10
819/819          7s 9ms/step -
loss: 0.0958 - mae: 0.2436 - val_loss: 0.1341 - val_mae: 0.2880
Epoch 7/10
819/819          8s 10ms/step -
loss: 0.0924 - mae: 0.2390 - val_loss: 0.1405 - val_mae: 0.2937
Epoch 8/10
819/819          7s 9ms/step -
loss: 0.0901 - mae: 0.2368 - val_loss: 0.1309 - val_mae: 0.2838
Epoch 9/10
819/819          8s 10ms/step -
loss: 0.0877 - mae: 0.2333 - val_loss: 0.1517 - val_mae: 0.3058
Epoch 10/10
819/819          8s 10ms/step -
loss: 0.0861 - mae: 0.2314 - val_loss: 0.1436 - val_mae: 0.2975
405/405          3s 6ms/step -
loss: 0.1476 - mae: 0.3039
Test MAE: 0.30
```

Plotting the results:

```
[50]: loss = history.history["mae"]
      val_loss = history.history["val_mae"]
      epochs = range(1, len(loss) + 1)
      plt.figure()
      plt.plot(epochs, loss, label="Training MAE")
      plt.plot(epochs, val_loss, label="Validation MAE")
      plt.legend()
```

```
[50]: <matplotlib.legend.Legend at 0x156aa83d0>
```



Trying 1D convolution method now. :

```
[54]: inputs = keras.Input(shape=(sequence_length, dfs.shape[-1]))
x = layers.Conv1D(8, 24, activation="relu")(inputs)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 12, activation="relu")(x)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 6, activation="relu")(x)
x = layers.GlobalAveragePooling1D()(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)
callbacks = [
    keras.callbacks.ModelCheckpoint("/Users/devmarwah/Documents/MSBA_
    ↪assignments/Advanced Machine Learning/Weather Forecasting/jena_conv.keras",
                                   save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_data=train_data, validation_data=validation_data,
                    epochs=10,
                    validation_data=Validation,
                    callbacks=callbacks)
```

```

model = keras.models.load_model("/Users/devmarwah/Documents/MSBA assignments/
↳Advanced Machine Learning/Weather Forecasting/jena_conv.keras")
print(f"Test MAE: {model.evaluate(Test)[1]:.2f}")

```

```

Epoch 1/10
819/819          14s 16ms/step -
loss: 0.2943 - mae: 0.4259 - val_loss: 0.2469 - val_mae: 0.3924
Epoch 2/10
819/819          13s 16ms/step -
loss: 0.1948 - mae: 0.3501 - val_loss: 0.2110 - val_mae: 0.3604
Epoch 3/10
819/819          14s 17ms/step -
loss: 0.1721 - mae: 0.3282 - val_loss: 0.2262 - val_mae: 0.3734
Epoch 4/10
819/819          14s 17ms/step -
loss: 0.1586 - mae: 0.3147 - val_loss: 0.2110 - val_mae: 0.3617
Epoch 5/10
819/819          15s 18ms/step -
loss: 0.1474 - mae: 0.3031 - val_loss: 0.2594 - val_mae: 0.3990
Epoch 6/10
819/819          15s 18ms/step -
loss: 0.1408 - mae: 0.2960 - val_loss: 0.1999 - val_mae: 0.3521
Epoch 7/10
819/819          15s 18ms/step -
loss: 0.1344 - mae: 0.2895 - val_loss: 0.1994 - val_mae: 0.3506
Epoch 8/10
819/819          15s 19ms/step -
loss: 0.1299 - mae: 0.2845 - val_loss: 0.2295 - val_mae: 0.3775
Epoch 9/10
819/819          15s 19ms/step -
loss: 0.1259 - mae: 0.2803 - val_loss: 0.2242 - val_mae: 0.3713
Epoch 10/10
819/819          15s 19ms/step -
loss: 0.1225 - mae: 0.2769 - val_loss: 0.2557 - val_mae: 0.3986
405/405          4s 9ms/step -
loss: 0.2342 - mae: 0.3832
Test MAE: 0.38

```

Plotting the results. :

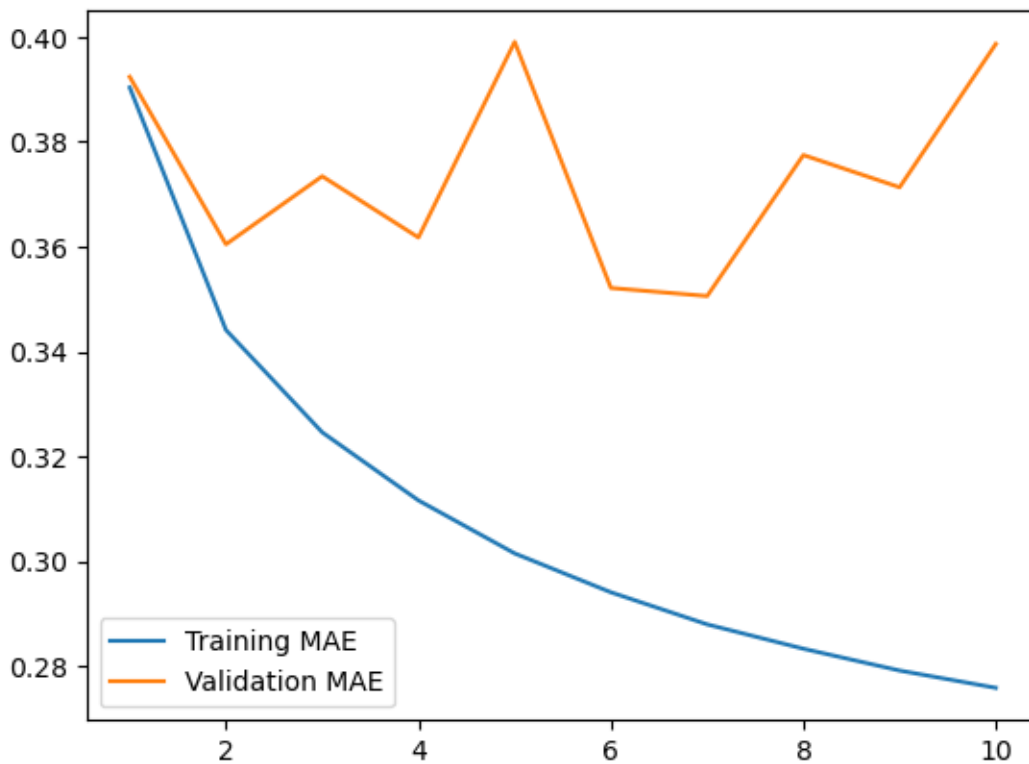
```

[55]: loss = history.history["mae"]
      val_loss = history.history["val_mae"]
      epochs = range(1, len(loss) + 1)
      plt.figure()
      plt.plot(epochs, loss, label="Training MAE")
      plt.plot(epochs, val_loss, label="Validation MAE")
      plt.legend()

```



[55]: <matplotlib.legend.Legend at 0x28e1d5810>



Constructing a simple recurrent neural network using LSTM model :

```
[62]: inputs = keras.Input(shape=(sequence_length,dfs.shape[-1]))
x = layers.LSTM(16)(inputs)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs,outputs)

callbacks = [
    keras.callbacks.ModelCheckpoint("Users/devmarwah/Documents/MSBA assignments/
↳Advanced Machine Learning/Weather Forecasting/jena_lstm.keras",
                                   save_best_only=True)
]

model.compile(optimizer="rmsprop",loss = "mse",metrics=["mae"])

history=model.fit(Train,
                  epochs=10,
                  validation_data=Validation,
                  callbacks=callbacks)
```

```
model=keras.models.load_model("Users/devmarwah/Documents/MSBA assignments/  
↳Advanced Machine Learning/Weather Forecasting/jena_lstm.keras")
```

```
Epoch 1/10  
819/819          22s 26ms/step -  
loss: 0.2250 - mae: 0.3494 - val_loss: 0.1148 - val_mae: 0.2631  
Epoch 2/10  
819/819          22s 26ms/step -  
loss: 0.1193 - mae: 0.2701 - val_loss: 0.1145 - val_mae: 0.2625  
Epoch 3/10  
819/819          22s 27ms/step -  
loss: 0.1099 - mae: 0.2598 - val_loss: 0.1171 - val_mae: 0.2652  
Epoch 4/10  
819/819          23s 28ms/step -  
loss: 0.1045 - mae: 0.2539 - val_loss: 0.1193 - val_mae: 0.2672  
Epoch 5/10  
819/819          22s 27ms/step -  
loss: 0.0999 - mae: 0.2485 - val_loss: 0.1247 - val_mae: 0.2740  
Epoch 6/10  
819/819          23s 28ms/step -  
loss: 0.0957 - mae: 0.2436 - val_loss: 0.1271 - val_mae: 0.2769  
Epoch 7/10  
819/819          23s 28ms/step -  
loss: 0.0922 - mae: 0.2390 - val_loss: 0.1278 - val_mae: 0.2776  
Epoch 8/10  
819/819          22s 27ms/step -  
loss: 0.0887 - mae: 0.2343 - val_loss: 0.1276 - val_mae: 0.2776  
Epoch 9/10  
819/819          23s 28ms/step -  
loss: 0.0855 - mae: 0.2299 - val_loss: 0.1366 - val_mae: 0.2863  
Epoch 10/10  
819/819          22s 27ms/step -  
loss: 0.0828 - mae: 0.2262 - val_loss: 0.1446 - val_mae: 0.2940
```


```
[65]: print("Test MAE: ",round(model.evaluate(Test)[1],2))
```

```
405/405          4s 10ms/step -  
loss: 0.1252 - mae: 0.2753  
Test MAE:  0.28
```

Plotting results of LSTM model :

```
[66]: loss = history.history['mae']  
val_loss = history.history["val_mae"]  
epochs = range(1, len(loss) + 1)  
plt.figure()  
plt.plot(epochs, loss,label="Training MAE")  
plt.plot(epochs,val_loss,label="Validation MAE")  
plt.legend()
```

[66]: <matplotlib.legend.Legend at 0x28fa48f90>



Hence, we can conclude that recurrent neural networks work best on a time - series problem

## 1.2 Improving Forecasting

Now we will use different model architectures to achieve the best possible accuracy in our model.

Constructing a RNN with only one layer of GRU but increased units to check initial accuracy.

```
[69]: inputs = keras.Input(shape=(sequence_length, dfs.shape[-1]))
x = layers.GRU(32, recurrent_dropout=0.25)(inputs)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)
callbacks = [
    keras.callbacks.ModelCheckpoint("/Users/devmarwah/Documents/MSBA_
    ↪assignments/Advanced Machine Learning/Weather Forecasting/jena_gru_dropout.
    ↪keras",
                                save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
```

```

history = model.fit(Train,
                    epochs=10,
                    validation_data=Validation,
                    callbacks=callbacks)
model = keras.models.load_model("Users/devmarwah/Documents/MSBA assignments/
↳Advanced Machine Learning/Weather Forecasting/jena_gru_dropout.keras")

```

```

Epoch 1/10
819/819          57s 68ms/step -
loss: 0.2557 - mae: 0.3812 - val_loss: 0.1187 - val_mae: 0.2693
Epoch 2/10
819/819          61s 75ms/step -
loss: 0.1643 - mae: 0.3157 - val_loss: 0.1137 - val_mae: 0.2630
Epoch 3/10
819/819          51s 62ms/step -
loss: 0.1602 - mae: 0.3113 - val_loss: 0.1144 - val_mae: 0.2640
Epoch 4/10
819/819          57s 70ms/step -
loss: 0.1569 - mae: 0.3085 - val_loss: 0.1117 - val_mae: 0.2601
Epoch 5/10
819/819          53s 65ms/step -
loss: 0.1528 - mae: 0.3043 - val_loss: 0.1126 - val_mae: 0.2614
Epoch 6/10
819/819          50s 61ms/step -
loss: 0.1497 - mae: 0.3016 - val_loss: 0.1118 - val_mae: 0.2606
Epoch 7/10
819/819          53s 65ms/step -
loss: 0.1454 - mae: 0.2972 - val_loss: 0.1137 - val_mae: 0.2627
Epoch 8/10
819/819          51s 62ms/step -
loss: 0.1421 - mae: 0.2942 - val_loss: 0.1172 - val_mae: 0.2672
Epoch 9/10
819/819          51s 62ms/step -
loss: 0.1398 - mae: 0.2912 - val_loss: 0.1121 - val_mae: 0.2604
Epoch 10/10
819/819          56s 68ms/step -
loss: 0.1365 - mae: 0.2883 - val_loss: 0.1147 - val_mae: 0.2649

```

```
[71]: model.evaluate(Test)
```

```

405/405          9s 21ms/step -
loss: 0.1277 - mae: 0.2792

```

```
[71]: [0.1279725730419159, 0.279487669467926]
```


```

[72]: loss = history.history['mae']
      val_loss = history.history["val_mae"]
      epochs = range(1, len(loss) + 1)

```

```
plt.figure()
plt.plot(epochs, loss,label="Training MAE")
plt.plot(epochs,val_loss,label="Validation MAE")
plt.legend()
```

[72]: <matplotlib.legend.Legend at 0x2d5158f90>



Now, we will take measures to improve our accuracy.

We are using following points to increase our model's efficiency :

- Adjusting the number of units in each recurrent layer in the stacked setup.
- Using `layer_lstm()` instead of `layer_gru()`.
- Using a combination of `1d_convnets` and RNN.

### Final Model Construction

```
[89]: inputs = keras.Input(shape=(sequence_length,dfs.shape[-1]))
# Using increased units and applying lstm instead of gru
x=layers.LSTM(32,recurrent_dropout=0.5,return_sequences=True)(inputs)
# Adding more layers to turn it into a stacked model
x=layers.LSTM(32,recurrent_dropout=0.5,return_sequences=True)(x)
# Applying 1D convolution
x = layers.Conv1D(8, 24, activation="relu")(x)
```

```

x = layers.MaxPooling1D(2)(x)
x = layers.GlobalAveragePooling1D()(x)
# We are also using dropout layer to regularize our results
x=layers.Dropout(0.5)(x)
outputs=layers.Dense(1)(x)
model=keras.Model(inputs,outputs)

callbacks = [
    keras.callbacks.ModelCheckpoint("/Users/devmarwah/Documents/MSBA_
    ↪assignments/Advanced Machine Learning/Weather Forecasting/jena_stacked_final.
    ↪keras",
                                save_best_only = True)
]

model.compile(optimizer="rmsprop",loss="mse",metrics=["mae"])

history=model.fit(Train,
                  epochs=10,
                  validation_data=Validation,
                  callbacks=callbacks)

```

```

Epoch 1/10
819/819          108s 131ms/step -
loss: 0.4698 - mae: 0.5320 - val_loss: 0.2734 - val_mae: 0.4102
Epoch 2/10
819/819          107s 131ms/step -
loss: 0.3609 - mae: 0.4599 - val_loss: 0.2372 - val_mae: 0.3828
Epoch 3/10
819/819          107s 130ms/step -
loss: 0.3285 - mae: 0.4324 - val_loss: 0.2543 - val_mae: 0.4015
Epoch 4/10
819/819          106s 130ms/step -
loss: 0.3057 - mae: 0.4149 - val_loss: 0.2474 - val_mae: 0.3969
Epoch 5/10
819/819          107s 130ms/step -
loss: 0.2961 - mae: 0.4056 - val_loss: 0.2600 - val_mae: 0.4028
Epoch 6/10
819/819          106s 130ms/step -
loss: 0.2894 - mae: 0.3987 - val_loss: 0.2493 - val_mae: 0.3951
Epoch 7/10
819/819          105s 128ms/step -
loss: 0.2827 - mae: 0.3929 - val_loss: 0.2592 - val_mae: 0.4054
Epoch 8/10
819/819          106s 129ms/step -
loss: 0.2792 - mae: 0.3894 - val_loss: 0.2539 - val_mae: 0.3993
Epoch 9/10
819/819          106s 129ms/step -
loss: 0.2722 - mae: 0.3841 - val_loss: 0.2698 - val_mae: 0.4143

```

```
Epoch 10/10  
819/819          107s 130ms/step -  
loss: 0.2702 - mae: 0.3807 - val_loss: 0.2670 - val_mae: 0.4088
```

```
[95]: model=keras.models.load_model("/Users/devmarwah/Documents/MSBA assignments/  
      ↪Advanced Machine Learning/Weather Forecasting/jena_stacked_final.keras")  
      model.evaluate(Test)[1]
```

```
405/405          17s 41ms/step -  
loss: 0.2792 - mae: 0.4163
```

```
[95]: 0.41434764862060547
```

Hence, our final model can predict temperature with a very low mean absolute error of 0.42